Machine Learning in Finance - Homework 5 – Due 10.08.2024

Textbook reading:
Chapters 5 starting page 129

1. **for a range of inputs from -infinity to +infinity, what are the output ranges for the elu, exponential, gelu, linear, and softmax activation functions in https://www.tensorflow.org/api_docs/python/tf/keras/activations**

   - ReLU:

     $$\max(x, 0)$$
     $$\in (0,\infty)$$

   - Exponential Linear Unit (ELU):

     $$\begin{matrix} \alpha(e^x - 1), & \text{if } x < 0 \\ x, & \text{if } x \geq 0 \end{matrix}$$
     $$\in (-\alpha, \infty)$$

   - Gaussian Error Linear Unit (GELU):

     $$\frac{1}{2}x\left(1 + \text{erf}\left(\frac{x}{\sqrt{2}}\right)\right) = x\Phi(x)$$
     $$\in (-0.17, \infty)$$

   - Linear (pass-through):

     $$x$$
     $$\in (-\infty, \infty)$$

   - Softmax:

     $$\frac{e^{x_i}}{\Sigma e^{x_i}}$$
     $$\in (0,1)$$

     It also sums to one.

2. **What is regularization and why does it help create better models?**

   Regularization adds a penalty term to the model's loss function, discouraging it from fitting the training data too closely by controlling the magnitude of model parameters.

   On the bias-variance trade off, regularization reduces variance by adding bias.

3. **Using matplotlib plot the swish function formula. How is it different from 'relu'? What advantages does it have over relu? What disadvantages does it have?**

   - Swish (SiLU):

$$\frac{x}{1+e^{-x}}$$

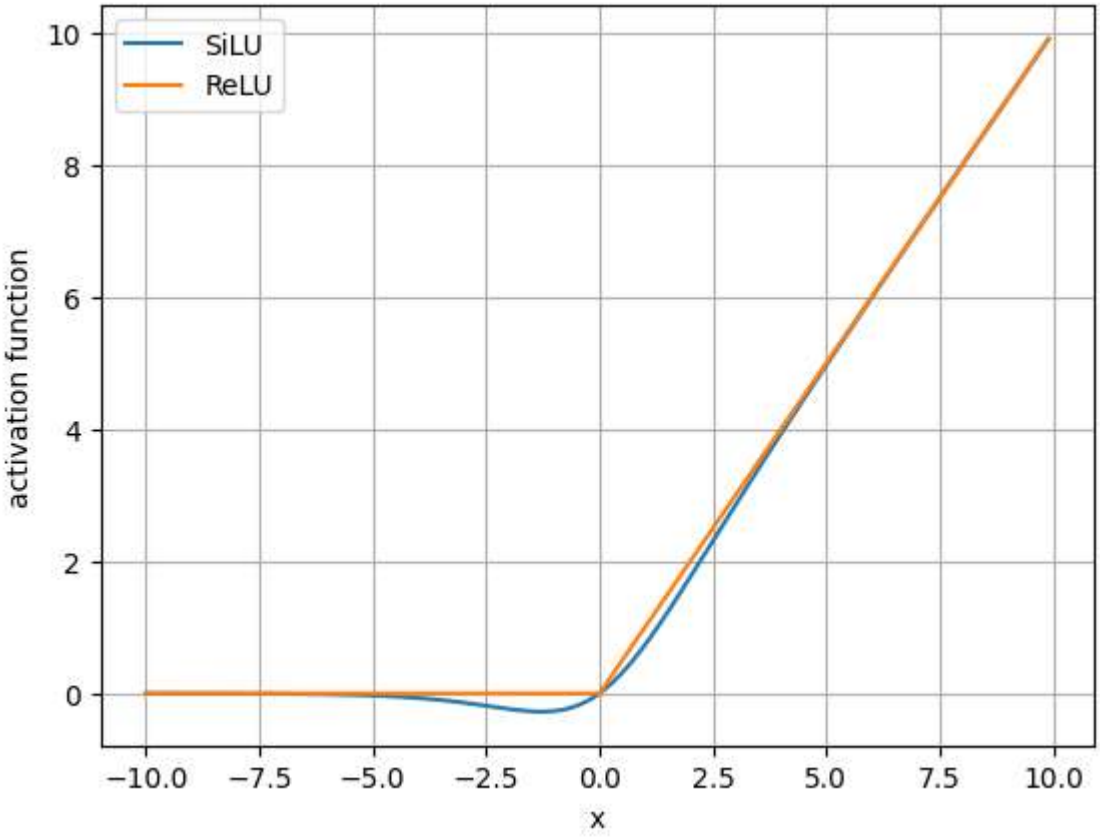

ReLU has a kink at x=0. SiLU is differentiable everywhere, which has an advantage over ReLU in gradient calculations. On the other hand, the gradients of SiLU is slightly more computational costly.

4. **In Section5.2_ANN.ipynb, replace 'relu' with 'swish' and replace the single layer model with a model containing two hidden layers, each containing 45 neurons.**

    a. **Plot the loss history and the accuracy history. How do they compare to the same model using 'relu' instead of 'swish'?**

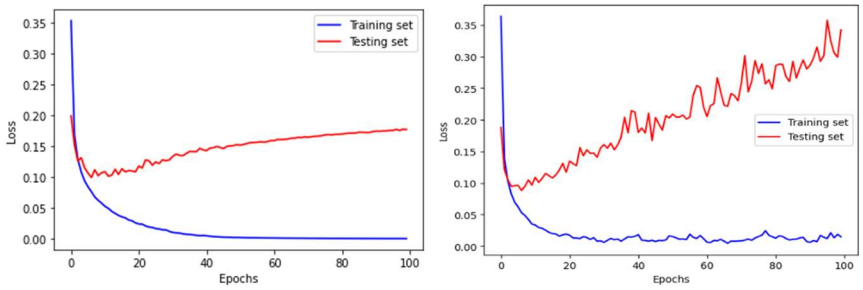

Figure 1. Loss history of ReLU(left) vs. SiLU(right)

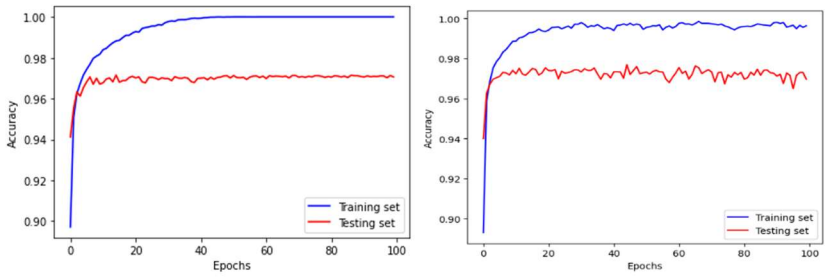

Figure 2. Accuracy history of ReLU(left) vs. SiLU(right)

By changing the model to 2 layers with SiLU activation function and 45 neurons each layer, the model seems to have overfitted the training dataset.

b. How many parameters are estimated in this model. Use the formulas to compute the number of parameters to verify that the input values match the actual number of parameters.

Input layer: 784

Hidden layer #1: 45

Hidden layer #2: 45

Output layer: 10

$$W_l = H_{l-1} \times H_l, \quad l = 1, 2, \ldots, L$$

$$B = H_1 + H_2 + \cdots + H_L$$

Total Parameters

= (784 x 45 + 45) + (45 x 45 + 45) + (45 x 10 + 10)

= 35325 + 2070 + 460 = 37855

Layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 45)	35,325
dense_7 (Dense)	(None, 45)	2,070
dense_8 (Dense)	(None, 10)	460

Total params: 37,855 (147.87 KB)
Trainable params: 37,855 (147.87 KB)
Non-trainable params: 0 (0.00 B)

c. **Which layers in the model have no parameters to fit?**

Input layer.