

temp

October 12, 2024



```
[19]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import os
# base_path = "/kaggle/input/playground-series-s4e10"
df_train = pd.read_csv(os.path.join('train.csv'))
df_test = pd.read_csv(os.path.join('test.csv'))
df_sample = pd.read_csv(os.path.join('sample_submission.csv'))
TARGET = 'loan_status'
```

Preprocessing

```
[20]: for column in df_train.columns:
    nan_count = df_train[column].isna().sum()
    print(f"Column '{column}' has {nan_count} NaN values.")
```

```
Column 'id' has 0 NaN values.
Column 'person_age' has 0 NaN values.
Column 'person_income' has 0 NaN values.
Column 'person_home_ownership' has 0 NaN values.
Column 'person_emp_length' has 0 NaN values.
Column 'loan_intent' has 0 NaN values.
Column 'loan_grade' has 0 NaN values.
Column 'loan_amnt' has 0 NaN values.
```

Column 'loan\_int\_rate' has 0 NaN values.  
 Column 'loan\_percent\_income' has 0 NaN values.  
 Column 'cb\_person\_default\_on\_file' has 0 NaN values.  
 Column 'cb\_person\_cred\_hist\_length' has 0 NaN values.  
 Column 'loan\_status' has 0 NaN values.

```
[21]: from sklearn.preprocessing import LabelEncoder

label_encoder = LabelEncoder()

for col in df_train.select_dtypes(include=['object']).columns:
    df_train[col] = label_encoder.fit_transform(df_train[col])

for col in df_test.select_dtypes(include=['object']).columns:
    df_test[col] = label_encoder.fit_transform(df_test[col])

print(df_train.head())
```

	id	person_age	person_income	person_home_ownership	person_emp_length	\
0	0	37	35000	3	0.0	
1	1	22	56000	2	6.0	
2	2	29	28800	2	8.0	
3	3	30	70000	3	14.0	
4	4	22	60000	3	2.0	

	loan_intent	loan_grade	loan_amnt	loan_int_rate	loan_percent_income	\
0	1	1	6000	11.49	0.17	
1	3	2	4000	13.35	0.07	
2	4	0	6000	8.90	0.21	
3	5	1	12000	11.11	0.17	
4	3	0	6000	6.92	0.10	

	cb_person_default_on_file	cb_person_cred_hist_length	loan_status
0	0	14	0
1	0	2	0
2	0	10	0
3	0	5	0
4	0	3	0

```
[22]: print(df_train.describe())
```

	id	person_age	person_income	person_home_ownership	\
count	58645.000000	58645.000000	5.864500e+04	58645.000000	
mean	29322.000000	27.550857	6.404617e+04	1.673578	
std	16929.497605	6.033216	3.793111e+04	1.452534	
min	0.000000	20.000000	4.200000e+03	0.000000	
25%	14661.000000	23.000000	4.200000e+04	0.000000	
50%	29322.000000	26.000000	5.800000e+04	3.000000	
75%	43983.000000	30.000000	7.560000e+04	3.000000	

max	58644.000000	123.000000	1.900000e+06	3.000000
-----	--------------	------------	--------------	----------

	person_emp_length	loan_intent	loan_grade	loan_amnt \
count	58645.000000	58645.000000	58645.000000	58645.000000
mean	4.701015	2.519430	1.066638	9217.556518
std	3.959784	1.722896	1.046181	5563.807384
min	0.000000	0.000000	0.000000	500.000000
25%	2.000000	1.000000	0.000000	5000.000000
50%	4.000000	3.000000	1.000000	8000.000000
75%	7.000000	4.000000	2.000000	12000.000000
max	123.000000	5.000000	6.000000	35000.000000


	loan_int_rate	loan_percent_income	cb_person_default_on_file \
count	58645.000000	58645.000000	58645.000000
mean	10.677874	0.159238	0.148384
std	3.034697	0.091692	0.355484
min	5.420000	0.000000	0.000000
25%	7.880000	0.090000	0.000000
50%	10.750000	0.140000	0.000000
75%	12.990000	0.210000	0.000000
max	23.220000	0.830000	1.000000

	cb_person_cred_hist_length	loan_status
count	58645.000000	58645.000000
mean	5.813556	0.142382
std	4.029196	0.349445
min	2.000000	0.000000
25%	3.000000	0.000000
50%	4.000000	0.000000
75%	8.000000	0.000000
max	30.000000	1.000000

```
[40]: X_train = df_train.drop(['id', 'loan_status'], axis=1).values
      # print(X_train[0,:])
      y_train = df_train['loan_status'].values

      X_test = df_test.values
```

```
[3.700e+01 3.500e+04 3.000e+00 0.000e+00 1.000e+00 1.000e+00 6.000e+03
 1.149e+01 1.700e-01 0.000e+00 1.400e+01]
```



```
[41]: from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
# from sklearn.metrics import roc_auc_score
# from sklearn.preprocessing import LabelEncoder
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_scaled = scaler.fit_transform(df_train.drop(['id', 'loan_status'], axis=1))
# Create the model
model = Sequential()
model.add(Dense(4, input_dim=X_scaled.shape[1], activation='relu')) # First hidden layer
model.add(Dense(4, activation='relu')) # Second hidden layer
model.add(Dense(1, activation='sigmoid')) # Output layer

# Compile the model
model.compile(loss='binary_crossentropy', optimizer='sgd')

# Fit the model
model.fit(X_scaled, y_train, epochs=100, batch_size=10, verbose=1)
```

Epoch 1/100

```
c:\Users\Big_Guppy\AppData\Local\Programs\Python\Python312\Lib\site-
packages\keras\src\layers\core\dense.py:87: UserWarning: Do not pass an
`input_shape`/`input_dim` argument to a layer. When using Sequential models,
prefer using an `Input(shape)` object as the first layer in the model instead.
  super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

5865/5865 18s 3ms/step -

loss: 0.3834

Epoch 2/100

5865/5865 15s 3ms/step -

loss: 0.2877

Epoch 3/100

5865/5865	12s 2ms/step -
loss: 0.2568	
Epoch 4/100	
5865/5865	19s 2ms/step -
loss: 0.2385	
Epoch 5/100	
5865/5865	24s 2ms/step -
loss: 0.2310	
Epoch 6/100	
5865/5865	20s 2ms/step -
loss: 0.2287	
Epoch 7/100	
5865/5865	22s 3ms/step -
loss: 0.2331	
Epoch 8/100	
5865/5865	28s 4ms/step -
loss: 0.2254	
Epoch 9/100	
5865/5865	32s 2ms/step -
loss: 0.2257	
Epoch 10/100	
5865/5865	25s 3ms/step -
loss: 0.2250	
Epoch 11/100	
5865/5865	20s 3ms/step -
loss: 0.2251	
Epoch 12/100	
5865/5865	17s 3ms/step -
loss: 0.2177	
Epoch 13/100	
5865/5865	15s 3ms/step -
loss: 0.2138	
Epoch 14/100	
5865/5865	14s 2ms/step -
loss: 0.2164	
Epoch 15/100	
5865/5865	27s 3ms/step -
loss: 0.2189	
Epoch 16/100	
5865/5865	17s 3ms/step -
loss: 0.2163	
Epoch 17/100	
5865/5865	22s 3ms/step -
loss: 0.2170	
Epoch 18/100	
5865/5865	18s 3ms/step -
loss: 0.2098	
Epoch 19/100	

5865/5865	19s 3ms/step -
loss: 0.2149	
Epoch 20/100	
5865/5865	17s 3ms/step -
loss: 0.2143	
Epoch 21/100	
5865/5865	17s 3ms/step -
loss: 0.2110	
Epoch 22/100	
5865/5865	18s 3ms/step -
loss: 0.2086	
Epoch 23/100	
5865/5865	17s 2ms/step -
loss: 0.2117	
Epoch 24/100	
5865/5865	13s 2ms/step -
loss: 0.2108	
Epoch 25/100	
5865/5865	22s 2ms/step -
loss: 0.2073	
Epoch 26/100	
5865/5865	16s 3ms/step -
loss: 0.2088	
Epoch 27/100	
5865/5865	12s 2ms/step -
loss: 0.2133	
Epoch 28/100	
5865/5865	12s 2ms/step -
loss: 0.2111	
Epoch 29/100	
5865/5865	13s 2ms/step -
loss: 0.2097	
Epoch 30/100	
5865/5865	15s 3ms/step -
loss: 0.2120	
Epoch 31/100	
5865/5865	12s 2ms/step -
loss: 0.2110	
Epoch 32/100	
5865/5865	29s 3ms/step -
loss: 0.2086	
Epoch 33/100	
5865/5865	28s 5ms/step -
loss: 0.2124	
Epoch 34/100	
5865/5865	44s 5ms/step -
loss: 0.2064	
Epoch 35/100	

5865/5865	32s 5ms/step -
loss: 0.2102	
Epoch 36/100	
5865/5865	36s 4ms/step -
loss: 0.2127	
Epoch 37/100	
5865/5865	44s 5ms/step -
loss: 0.2111	
Epoch 38/100	
5865/5865	46s 6ms/step -
loss: 0.2094	
Epoch 39/100	
5865/5865	16s 3ms/step -
loss: 0.2092	
Epoch 40/100	
5865/5865	20s 3ms/step -
loss: 0.2133	
Epoch 41/100	
5865/5865	18s 3ms/step -
loss: 0.2066	
Epoch 42/100	
5865/5865	17s 3ms/step -
loss: 0.2134	
Epoch 43/100	
5865/5865	16s 3ms/step -
loss: 0.2119	
Epoch 44/100	
5865/5865	25s 3ms/step -
loss: 0.2134	
Epoch 45/100	
5865/5865	24s 4ms/step -
loss: 0.2067	
Epoch 46/100	
5865/5865	17s 3ms/step -
loss: 0.2081	
Epoch 47/100	
5865/5865	21s 4ms/step -
loss: 0.2103	
Epoch 48/100	
5865/5865	22s 4ms/step -
loss: 0.2084	
Epoch 49/100	
5865/5865	38s 3ms/step -
loss: 0.2094	
Epoch 50/100	
5865/5865	30s 5ms/step -
loss: 0.2078	
Epoch 51/100	

5865/5865	27s 5ms/step -
loss: 0.2081	
Epoch 52/100	
5865/5865	46s 5ms/step -
loss: 0.2089	
Epoch 53/100	
5865/5865	35s 4ms/step -
loss: 0.2049	
Epoch 54/100	
5865/5865	46s 5ms/step -
loss: 0.2076	
Epoch 55/100	
5865/5865	38s 4ms/step -
loss: 0.2092	
Epoch 56/100	
5865/5865	25s 4ms/step -
loss: 0.2055	
Epoch 57/100	
5865/5865	24s 4ms/step -
loss: 0.2042	
Epoch 58/100	
5865/5865	46s 5ms/step -
loss: 0.2059	
Epoch 59/100	
5865/5865	40s 4ms/step -
loss: 0.2089	
Epoch 60/100	
5865/5865	37s 4ms/step -
loss: 0.2082	
Epoch 61/100	
5865/5865	54s 6ms/step -
loss: 0.2100	
Epoch 62/100	
5865/5865	31s 4ms/step -
loss: 0.2061	
Epoch 63/100	
5865/5865	45s 5ms/step -
loss: 0.2060	
Epoch 64/100	
5865/5865	24s 4ms/step -
loss: 0.2118	
Epoch 65/100	
5865/5865	40s 7ms/step -
loss: 0.2086	
Epoch 66/100	
5865/5865	52s 8ms/step -
loss: 0.2117	
Epoch 67/100	



5865/5865	61s 5ms/step -
loss: 0.2146	
Epoch 68/100	
5865/5865	23s 4ms/step -
loss: 0.2100	
Epoch 69/100	
5865/5865	23s 4ms/step -
loss: 0.2070	
Epoch 70/100	
5865/5865	40s 4ms/step -
loss: 0.2049	
Epoch 71/100	
5865/5865	22s 4ms/step -
loss: 0.2079	
Epoch 72/100	
5865/5865	23s 4ms/step -
loss: 0.2085	
Epoch 73/100	
5865/5865	22s 4ms/step -
loss: 0.2086	
Epoch 74/100	
5865/5865	40s 3ms/step -
loss: 0.2058	
Epoch 75/100	
5865/5865	22s 4ms/step -
loss: 0.2094	
Epoch 76/100	
5865/5865	22s 4ms/step -
loss: 0.2060	
Epoch 77/100	
5865/5865	40s 3ms/step -
loss: 0.2112	
Epoch 78/100	
5865/5865	20s 3ms/step -
loss: 0.2063	
Epoch 79/100	
5865/5865	19s 3ms/step -
loss: 0.2072	
Epoch 80/100	
5865/5865	21s 3ms/step -
loss: 0.2099	
Epoch 81/100	
5865/5865	19s 3ms/step -
loss: 0.2059	
Epoch 82/100	
5865/5865	19s 3ms/step -
loss: 0.2097	
Epoch 83/100	

5865/5865	31s 5ms/step -
loss: 0.2091	
Epoch 84/100	
5865/5865	38s 4ms/step -
loss: 0.2108	
Epoch 85/100	
5865/5865	24s 4ms/step -
loss: 0.2082	
Epoch 86/100	
5865/5865	27s 4ms/step -
loss: 0.2072	
Epoch 87/100	
5865/5865	39s 4ms/step -
loss: 0.2084	
Epoch 88/100	
5865/5865	18s 3ms/step -
loss: 0.2123	
Epoch 89/100	
5865/5865	16s 2ms/step -
loss: 0.2051	
Epoch 90/100	
5865/5865	18s 3ms/step -
loss: 0.2058	
Epoch 91/100	
5865/5865	19s 3ms/step -
loss: 0.2075	
Epoch 92/100	
5865/5865	20s 3ms/step -
loss: 0.2075	
Epoch 93/100	
5865/5865	17s 3ms/step -
loss: 0.2067	
Epoch 94/100	
5865/5865	17s 3ms/step -
loss: 0.2107	
Epoch 95/100	
5865/5865	17s 3ms/step -
loss: 0.2105	
Epoch 96/100	
5865/5865	22s 3ms/step -
loss: 0.2093	
Epoch 97/100	
5865/5865	17s 3ms/step -
loss: 0.2030	
Epoch 98/100	
5865/5865	17s 3ms/step -
loss: 0.2064	
Epoch 99/100	

```
5865/5865          23s 3ms/step -
loss: 0.2066
Epoch 100/100
5865/5865          20s 3ms/step -
loss: 0.2093
```

```
[41]: <keras.src.callbacks.history.History at 0x1500617b110>
```

```
[30]: print(df_test['id'])
```

```
0      58645
1      58646
2      58647
3      58648
4      58649
...
39093   97738
39094   97739
39095   97740
39096   97741
39097   97742
Name: id, Length: 39098, dtype: int64
```

```
[42]: # Assuming 'id' is the first column in the NumPy array
X_test_ids = df_test['id'] # Get the 'id' values (first column)
X_test_scaled = scaler.transform(df_test.drop('id', axis=1)) # Scale the test_
↳ set features

# Make predictions
y_test_pred_probs = model.predict(X_test_scaled).flatten()

# Prepare the submission DataFrame
submission_df = pd.DataFrame({
    'id': X_test_ids,
    'loan_status': y_test_pred_probs
})

# Save to CSV
submission_df.to_csv('submission.csv', index=False)
```

```
1222/1222          4s 3ms/step
```